

Task 1:

Most of this has been implemented in mdp.py. The MDP class has `Q_from_V`, `pi_from_Q`, and `V_from_pi` functions that are used to convert from `pi` to `V` or the other way around. All these are implemented straightforwardly from the definitions taught.

For `value_iteration`, we apply the bellman optimality operator (in vectorized form) till consecutive `V` and `V'` have rms difference less than `epsilon = 1e-10`. There is also a cap on the number of iterations, just in case.

For `howards_policy_iteration`, we set `pi_ = q_pi.argmax(axis=1)`. So if there are actions that have better `Q` value, we choose the one with the highest `Q`. Otherwise, the current action has the best `Q`, and hence `argmax` chooses that itself. We keep repeating until `pi` and `pi_` are the same. Tie breaking is the default followed by `numpy.argmax`, which is the action with the lowest index among those that are tied.

For `linear_programming`, I created `S` variables for each `V(S_i)`, and then created constraints to be added to the problem iteratively.

Task 2:

I create two extra states. The loss/draw terminal state, and the win terminal state. These states are the largest two indices and are mentioned as end states (but not used, since planner doesn't really use the terminal states). The transitions are normally transcribed. For each state action pair, the following is done:

if action is illegal, create transition to loss-state with -100 reward.

else, the state that the opponent will see because of this state + action, is created. Now based on the opponent's policy: 1. If game continues, the transition is printed with corresponding probability. 2. If game ends, this actions probability is added to one of `win_prob` / `draw_prob` which accumulate the win and draw probabilities. After accounting for all actions of the opponent, the cumulative win and draw transitions are printed.

in case the state created for the opponent is not in its policy, this implies the game ended by state + action. This implies the game has ended in a loss/draw and hence the appropriate transition is printed.

Task 3:

Yes, the policy always converges. It converges to an "ultimate policy", which is one which will win in all situations where it is possible to guarantee a win against all opponent policies. This implies that player 2 always wins from the blank board state too, no matter how player 1 were to play. These policies are such that they complement each other. So an ultimate policy for `p1` is an optimal policy when `p2` follows its ultimate policy, and the other way round too.

Proof: Let us say that we start with `p2's` policy `p20`, where `pij` is policy of player `i` in step `j`. This gives us `p11` and that gives us `p21`. We know that the Value functions of these players are complementary in a way. So if `p1` has a high value for a state, `p2` will have a low value for the state that follows from this state and the action that `p1` chooses. `V` of `p20` under `p11` is lower than `V` of `p21` under `p11`. This is because `p21` is the optimal policy against `p11` and hence has the highest `V` for each state. Now `p12` has to be such that `V` of `p12` under `p21` is higher than that of `p11` under `p21` (by the same argument). Hence value of `p21` under `p12` is lower than its value under `p11` (since values are complementary).

Therefore $V_{p20_p11} \leq V_{p21_p11} \geq V_{p21_p12}$.

Hence `p2's` policy's value first increases then decreases, and this keeps on happening, and can only stop when it becomes ultimate.

Similarly, $V_{p10_p21} \leq V_{p11_p21} \geq V_{p11_p22}$ implying $V_{p21_p10} \geq V_{p21_p11} \leq V_{p22_p11}$.

Therefore $V_{p20_p11} \leq V_{p22_p11}$ which means `V` of `p2` must increase in net after the increase and decrease, and this guarantees convergence to the ultimate policy.